

Search in Complex Environments

Presented by Yasin Ceran

September 30, 2024

1 Hill-climbing

2 Simulated Annealing

3 Genetic Algorithms

4 Continuous State Spaces

Iterative Improvement Algorithms

Module4 focused on problems in:

- Fully observable
- Deterministic
- Static
- Known environments

where the solution is a **sequence of actions**.

In Module 5 we relax these constraints:

- **Finding good states** without considering the path (discrete or continuous paths)
- **Relaxing determinism:**
 - Agent needs a **conditional plan**, e.g., stop if red, go if green.
- **Partial observability:**
 - Agent tracks **possible states** it might be in.
- **Online search in unknown spaces:**
 - Agent must **learn as it goes**.

Local Search Algorithms

Key Characteristics of Local Search Algorithms:

- Search from a **start state** to neighboring states.
- **Do not track** the path or previously reached states.
- Not systematic—might miss parts of the search space where a solution exists.

Advantages of Local Search:

- Use very **little memory**.
- Can often find **reasonable solutions** in large or infinite state spaces.

Local search algorithms can also solve **optimization problems**, in which the aim is to find the best state according to an **objective function**.

Hill-climbing (or gradient ascent/descent)

“Like climbing Everest in thick fog with amnesia”

function Hill-Climbing(*problem*) **returns** a state that is a local maximum

inputs: *problem*, a problem

local variables: *current*, a node
neighbor, a node

current \leftarrow Make-Node(Initial-State[*problem*])

loop do

neighbor \leftarrow a highest-valued successor of *current*


if Value[neighbor] \leq Value[current] **then return** State[*current*]

current \leftarrow *neighbor*

end

Hill-climbing contd.

Useful to consider **state space landscape**




Random-restart hill climbing overcomes local maxima—trivially complete

Random sideways moves 😊 escape from shoulders 🙄 loop on flat maxima

Example: n -queens


Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Move a queen to reduce number of conflicts



Almost always solves n -queens problems almost instantaneously for very large n , e.g., $n = 1$ million

Example: 8-queens



(a)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	13	16	13	16	16
17	14	17	15	14	16	16	16
17	16	18	15	14	15	16	16
18	14	15	15	14	16	16	16
14	14	13	17	12	14	12	18

(b)

The board shows the value of h for each possible successor obtained by moving a queen within its column.

Simulated Annealing

Idea: escape local maxima by allowing some “bad” moves
but gradually decrease their size and frequency

function *Simulated-Annealing*(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

local variables: *current*, a node

next, a node

T, a “temperature” controlling prob. of downward

steps

current \leftarrow Make-Node(Initial-State[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ Value[*next*] – Value[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{-\Delta E / T}$

Properties of Simulated Annealing

At fixed “temperature” T , state occupation probability reaches Boltzman distribution

$$p(x) = \alpha e^{\frac{E(x)}{kT}}$$

T decreased slowly enough \implies always reach best state x^*
 because $e^{\frac{E(x^*)}{kT}} / e^{\frac{E(x)}{kT}} = e^{\frac{E(x^*) - E(x)}{kT}} \gg 1$ for small T

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process modelling

Widely used in VLSI layout, airline scheduling, etc.

Local Beam Search

Idea: keep k states instead of 1; choose top k of all their successors

Not the same as k searches run in parallel!

Searches that find good states recruit other searches to join them



Problem: quite often, all k states end up on same local hill

Idea: choose k successors randomly ([Stochastic Beam Search](#)), biased towards good ones

Observe the close analogy to natural selection!

Genetic Algorithms

= stochastic local beam search + generate successors from **pairs** of states



Genetic Algorithms-II

function Genetic Algorithm(*population*, *Fitness-Fn*) **returns** an individual

inputs: *population*, a set of individuals

Fitness-Fn, a function that measures the fitness of an individual

repeat

new_population \leftarrow empty set

for $i=1$ **to** Size[*population*] **do**

$x \leftarrow$ Random-Selection(*population*, [Fitness-Fn])

$y \leftarrow$ Random-Selection(*population*, [Fitness-Fn])

child \leftarrow Reproduce(x , y)

if (small random probability) **then** *child* \leftarrow Mutate(*child*)

add *child* to *new_population*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to Fitness-Fn

Genetic Algorithms-III

```
function REPRODUCE( x,y) returns an individual
  inputs: x , y , parent, individuals,

  n ← LENGTH(x)
  c ← random number from 1 to n
  return APPEND(SUBSTRING(x,1,c),SUBSTRING(y,c +1,n)
```

Continuous State Spaces

Suppose we want to site three airports in Romania:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space,
e.g., empirical gradient considers $\pm\delta$ change in each coordinate

Gradient methods compute

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Sometimes can solve for $\nabla f(\mathbf{x}) = 0$ exactly (e.g., with one city).

Summary

- Hill-climbing is simply a loop that continually moves in the direction of increasing value
- With simulated annealing, we can escape local maxima by allowing some “bad” moves
- Local beam search allows us to keep track of multiple states instead of 1, thus choose k successors that are biased towards the good ones
- Continuous state spaces can offer us search algorithms that can be used for real life continuous environments.