

# Adversarial Games

Presented by Yasin Ceran

October 29, 2024

1 Games

2 Optimal Decisions in Games

3 Heuristic ABTS

4 MCTS

5 Stochastic Games

# Game theory

There are three stances we can take towards multi-agent environments:

- When there is a very large number of agents, consider them in the aggregate as an **economy**, e.g., predicting that increasing demand will cause prices to rise
- We would consider adversarial agents as just a part of the environment, like rain sometimes falls and sometimes does not. In this case, we miss the idea that adversaries are actively trying to defeat us.
- Explicitly model the adversarial agents with the techniques of adversarial game-tree search

# Game Theory

**Definition:** Game theory is the study of mathematical models of strategic interaction between rational decision-makers.

- Focuses on decision-making in competitive environments.
- Players aim to maximize their payoff by choosing optimal strategies.
- Types of games:
  - Zero-sum games: One player's gain is another player's loss.
  - Non-zero-sum games: Both players can win or lose together.
- Common solution concepts:
  - Nash equilibrium
  - Dominant strategies
  - Minimax strategy

# Nash Equilibrium

**Definition:** A Nash equilibrium is a situation in a game where no player can benefit by changing their strategy while the other players keep their strategies unchanged.

- At Nash equilibrium, all players are playing optimally given the strategies of others.
- No player has an incentive to deviate from their strategy.
- A game can have one, multiple, or no Nash equilibria.

## Example: Prisoner's Dilemma

- The dominant strategy for both prisoners is to "confess," resulting in a Nash equilibrium where both confess.
- Even though mutual silence would have been better for both, the rational choice leads them to confess.

# Dominant Strategies

**Definition:** A strategy is dominant if, regardless of what the other players do, the strategy earns a player a higher payoff than any other.

- A dominant strategy is the best strategy for a player, no matter what the opponents decide.
- If all players in a game have a dominant strategy, the outcome of the game is determined.
- Not all games have dominant strategies for all players.

## Key Points:

- A player with a dominant strategy will always choose it.
- The existence of a dominant strategy simplifies the decision-making process for that player.

## Example:

- In the Prisoner's Dilemma, the dominant strategy for each player is to "confess," as it leads to a better outcome for them regardless of the opponent's choice.

# Optimal Decisions in Games

**Goal:** Find the best possible move in a game given the current state.

- **Minimax Algorithm:** A decision rule for minimizing the possible loss in a worst-case scenario.
  - Each player selects the move that maximizes their payoff, assuming the opponent will minimize it.
  - Typically applied in two-player, zero-sum games.
- **Minimax Principle:** Choose the action that maximizes the minimum gain (for the max player) or minimizes the maximum loss (for the min player).
- **Optimality:** The minimax algorithm leads to optimal decisions in deterministic games where all information is available.

# Games in General

- In **multiagent environments**, each agent considers the actions of the other agents
- In **competitive** environments, agents' goals are in conflict. This competition results in **adversarial search**
- A game of **perfect information** happens in a deterministic and fully observable environments
- In a game of **imperfect information**, agents don't see the actions of the other agents but they may know the possibilities
- **Pruning** allows us to ignore the portions of the search tree that make no difference to the final choice.
- The heuristic **evaluation functions** allow us to approximate the true utility of a state without doing a complete search.


	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker, scrabble nuclear war



# Two-Player Zero-Sum Games

- **zero-sum** means what is good for one player is just as bad for the other
- **move** is a synonym for “action”, and **position** is a synonym for “state”
- We call our two players ‘MAX’ and ‘MIN’
- $S_0$ : The **initial state**, which specifies how the game is set up at the start
- $TO - MOVE(s)$ : The player whose turn it is to move in state  $s$
- $ACTIONS(s)$ : The set of legal moves in state  $s$
- $RESULT(s, a)$ : The **transition model** defines the state resulting from taking action  $a$  in state  $s$
- $IS - TERMINAL(s)$ : A terminal test, which is true when the game is over and false otherwise
- $UTILITY(s, p)$ : A **utility function** (objective function or payoff function) defines the final numeric value to player  $p$  when the game ends in terminal state  $s$
- **game tree** is a search tree that follows every sequence of moves all the way to a terminal state

# Example: Game tree (2-player, deterministic, turns)




# Minimax

- Perfect play for deterministic, perfect-information games
- MAX wants to find a sequence of actions to a win, but MIN has something to say about it; MAX's strategy must be a conditional plan
- Idea: choose move to position with highest

**minimax value** = best achievable payoff against best play

E.g., 2-ply game:



# Minimax Algorithm

**function** MINIMAX\_SEARCH(*game, state*) *returns an action*

*inputs:* *state*, current state in game

player  $\leftarrow$  game.TO\_MOVE(*state*)

*value, move*  $\leftarrow$  MAX\_VALUE(*game, state*)

**return** *move*

**function** MAX\_VALUE(*game, state*) *returns a (utility, move) pair*

**if** game.IS\_TERMINAL(*state*) **then return** game.UTILITY(*(state, player)*, null)

*v*  $\leftarrow -\infty$

**for each** *a* **in** game.ACTIONS(*state*) **do**

*v2, a2*  $\leftarrow$  MIN\_VALUE(*game, game.RESULT(state, a)*)

**if** *v2* > *v* **then**

*v, move*  $\leftarrow$  *v2, a*

**return** *v, move*

**function** MIN\_VALUE(*game, state*) *returns a (utility, move) pair*

**if** game.IS\_TERMINAL(*state*) **then return** game.UTILITY(*(state, player)*, null)

*v*  $\leftarrow +\infty$

**for each** *a* **in** game.ACTIONS(*state*) **do**

*v2, a2*  $\leftarrow$  MAX\_VALUE(*game, game.RESULT(state, a)*)







**if** *v2* < *v* **then**


*v, move*  $\leftarrow$  *v2, a*

**return** *v, move*

# $\alpha$ - $\beta$ Pruning

- The number of game states is exponential in the depth of the tree.
- Pruning** large parts of the tree that make no difference to the outcome will reduce the number of states





- If  $m$  or  $m'$  is better than  $n$  for Player, we will never get to  $n$  in play.
- $\alpha$ : The value of the best (highest value) choice we have found so far at any choice along the path for MAX. Think:  $\alpha = \text{'at\_least'}$
- $\beta$ : The value of the best (lowest value) choice we have found so far at any choice along the path for MIN. Think:  $\beta = \text{'at\_most'}$

# Alpha-Beta Algorithm

```

function ALPHA_BETA_SEARCH(game, state) returns an action
  inputs: state, current state in game

  player  $\leftarrow$  game.TO_MOVE(state)
  value, move  $\leftarrow$  MAX_VALUE(game, state,  $-\infty$ ,  $+\infty$ )
  return move

```

---

```

function MAX_VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
  if game.IS_TERMINAL(state) then return game.UTILITY(state,
    player), null)
  v  $\leftarrow$   $-\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MIN_VALUE(game, game.RESULT(state, a),  $\alpha$ ,  $\beta$ )
    if v2 > v then
      v, move  $\leftarrow$  v2, a
       $\alpha \leftarrow$  (MAX)( $\alpha$ , v)
    if v  $\geq$   $\beta$  then return v, move
  return v, move

```

---

```

function MIN_VALUE(game, state,  $\alpha$ ,  $\beta$ ) returns a (utility, move) pair
  if game.IS_TERMINAL(state) then return game.UTILITY(state,
    player), null)
  v  $\leftarrow$   $+\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2  $\leftarrow$  MAX_VALUE(game, game.RESULT(state, a),  $\alpha$ ,  $\beta$ )
    if v2 < v then
      v, move  $\leftarrow$  v2, a
       $\beta \leftarrow$  (MIN)( $\beta$ , v)
    if v  $\leq$   $\alpha$  then return v, move
  return v, move


```

# Evaluation Functions

- A heuristic evaluation function  $EVAL(s, p)$  returns an **estimate** of the expected utility of state  $s$  to player  $p$
- For terminal states,  $EVAL(s, p) = UTILITY(s, p)$
- For non-terminal states,  
 $UTILITY(loss, p) \leq EVAL(s, p) \leq UTILITY(win, p)$
- The computation must not take too long
- The evaluation function should be strongly correlated with the actual chances of winning
- Most evaluation functions work by calculating various **features** of the state
- Most evaluation functions compute separate numerical contributions from each feature and then **combine** them to find the total value




# Evaluation Functions-Chess



Black to move

White slightly better



White to move

Black winning

For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

# Monte Carlo Tree Search (MCTS)

- The basic **MCTS** strategy does not use a heuristic evaluation function
- The value of a state is estimated as the average utility over a number of **simulations** of complete games starting from the state
- A simulation chooses moves first for one player, than for the other, repeating until a terminal position is reached
- To get useful information from the simulation, we need **playout policy** that biases the moves towards the good ones
- **Selection** is starting at the root of the search tree and choosing a move guided by the selection policy
- **Expansion** is growing the search tree by a new child of the selected node
- **Simulation** is performing a playout from the newly generated child node, choosing move for both players according to playout policy
- **Back-propagation** is using the result of the simulation to update all the search tree nodes going up to the root


# MCTS Algorithm

```
function MONTE_CARLO_TREE_SEARCH(state) returns an action  
  
  tree ← NODE(state)  
  while IS_TIME_REMAINING() do  
    leaf ← SELECT(tree)  
    child ← EXPAND(leaf)  
    result ← SIMULATE(child)  
    function BACK_PROPOGATE(result, child) returns  
  return the move in ACTION(state) whose node has the highest  
  number of playouts
```

# Stochastic Games in General

In nondeterministic games, chance introduced by dice, card-shuffling

Simplified example with coin-flipping:



# Algorithm for Stochastic Games

We can only calculate **expected value** of a position: the average over all possible outcomes of the chance nodes

Expectiminimax gives perfect play

Just like Minimax, except we must also handle chance nodes:

```

EXPECTIMINIMAX( $s$ )=
  if IS_TERMINAL( $s$ )
    then UTILITY( $s$ , MAX)
  if TO_MOVE( $s$ )=MAX
    then  $\max_a$  EXPECTIMINIMAX(RESULT( $s$ ,  $a$ )))
  if TO_MOVE( $s$ )=MIN
    then  $\min_a$  EXPECTIMINIMAX(RESULT( $s$ ,  $a$ )))
  if TO_MOVE( $s$ )=CHANCE
    then  $\sum_r P(r)$  EXPECTIMINIMAX(RESULT( $s$ ,  $r$ )))
  
```

# Summary

Games are fun to work on! (and dangerous)

They illustrate several important points about AI

- Game theory and types of games (zero-sum, non-zero-sum)
- Optimal decisions in games using the Minimax algorithm
- Heuristic improvements with Alpha-Beta pruning
- Monte Carlo Tree Search for large and complex search spaces
- Stochastic games and decision-making under uncertainty